

```
#!/usr/bin/env python3
# -*- coding: utf-8 -*-
"""
```

Created on Sat Apr 27 21:59:47 2024

```
@author: ananyaashahi
"""
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix, classification_report
```

```
#loading the dataset
df=pd.read_csv("Application_Data 2.csv")
```

```
###
```

```
#printing the information of the dataset
```

```
print(df.info())
```

```
#   Column                Non-Null Count  Dtype
# ---  -
# 0   Applicant_ID        25128 non-null    int64
# 1   Applicant_Gender    25128 non-null    object
# 2   Owned_Car           25128 non-null    int64
# 3   Owned_Realty        25128 non-null    int64
# 4   Total_Children      25128 non-null    int64
# 5   Total_Income        25128 non-null    int64
# 6   Income_Type         25128 non-null    object
# 7   Education_Type      25128 non-null    object
# 8   Family_Status       25128 non-null    object
# 9   Housing_Type        25128 non-null    object
# 10  Owned_Mobile_Phone  25128 non-null    int64
# 11  Owned_Work_Phone    25128 non-null    int64
# 12  Owned_Phone         25128 non-null    int64
# 13  Owned_Email         25128 non-null    int64
# 14  Job_Title           25128 non-null    object
# 15  Total_Family_Members 25128 non-null    int64
# 16  Applicant_Age       25128 non-null    int64
# 17  Years_of_Working    25128 non-null    int64
# 18  Total_Bad_Debt      25128 non-null    int64
# 19  Total_Good_Debt     25128 non-null    int64
# 20  Status              25128 non-null    int64
# dtypes: int64(15), object(6)
```

```
# Categorical features are - ['Applicant_Gender', 'Income_Type', 'Education_Type',
#                             'Family_Status', 'Housing_Type', 'Job_Title']
```

```
###
```

```
print(df.shape)
```

```

#Our dataset has 25128 rows and 21 columns

###

#counting missing values per row

row_nan_count = df.isna().sum(axis=1)
print(row_nan_count)

# NO missing values
###

#checking for duplicate values
print("Number of duplicate rows", df.duplicated().sum())

# NO duplicate values
###

print(df.head())

###
#Checking the uniques values in the Job_Title
unique_job_titles = df['Job_Title'].unique()

#Counting occurrences of each unique value in the 'Job_Title' column
job_titles_counts = df['Job_Title'].value_counts()

# Display unique values
print("Unique Job Titles:")
print(unique_job_titles)

# Display counts of each employment type
print("\nJob Title Counts:")
print(job_titles_counts)


#Job Title Counts:
#Laborers 6211
#Core staff 3591
#Sales staff 3485
#Managers 3012
#Drivers 2135
#High skill tech staff 1383
#Accountants 1241
#Medicine staff 1207
#Cooking staff 655
#Security staff 592
#Cleaning staff 549
#Private service staff 344
#Low-skill Laborers 175
#Waiters/barmen staff 173
#Secretaries 151
#HR staff 85
#Realty agents 79

```

```
#IT staff
```

60

```
###
```

```
# Removing the trailing space in the column Job_Title
# Noticed we had trailing space in that column, which was preventing the conversion
of grouping of Job_Titles
```

```
df['Job_Title'] = df['Job_Title'].str.strip()
```

```
###
```

```
job_title_mapping = {
    'Professional Roles': ['Managers', 'High skill tech staff', 'Accountants',
    'Medicine staff'],
    'Service and Support Roles': ['Sales staff', 'Cooking staff', 'Secretaries',
    'HR staff',
                                'Realty agents', 'Private service staff',
    'Waiters/barmen staff'],
    'Skilled Labor Roles': ['Drivers', 'Security staff'],
    'Unskilled Labor Roles': ['Laborers', 'Core staff', 'Cleaning staff', 'Low-
skill Laborers']
}
```

```
# Function to map job titles to broader categories
```

```
def map_job_title_to_category(job_title):
    for category, titles in job_title_mapping.items():
        if job_title in titles:
            return category
    return 'Other' # Default category for job titles not in the mapping
```

```
# Applying the mapping function to create a new column 'Job_Category'
```

```
df['Job_Category'] = df['Job_Title'].apply(map_job_title_to_category)
```

```
print(df.head())
```

```
# Validating if the Job_Category column contains the expected categories and their
respective counts.
```

```
job_category_counts = df['Job_Category'].value_counts()
```

```
print(job_category_counts)
```

```
#Dropping the Job-Title feature
```

```
df.drop('Job_Title', axis=1, inplace=True)
```

```
print(df.head())
```

```
#   Applicant_ID Applicant_Gender  ... Status      Job_Category
#0      5008806             M      ...     1  Skilled Labor Roles
#1      5008808             F      ...     1  Service and Support Roles
#2      5008809             F      ...     1  Service and Support Roles
#3      5008810             F      ...     1  Service and Support Roles
#4      5008811             F      ...     1  Service and Support Roles
```

```
#[5 rows x 21 columns]
```

```
###
```

```
# Checking the unique values of
'Owned_Mobile_Phone', Owned_Mobile_Phone, 'Owned_Phone', 'Owned_Email'
```

```
# Check unique values and counts for 'Owned_Mobile_Phone'
print("Unique values and counts for 'Owned_Mobile_Phone':")
print(df['Owned_Mobile_Phone'].value_counts())
print("\n")
```

```
# Check unique values and counts for 'Owned_Work_Phone'
print("Unique values and counts for 'Owned_Work_Phone':")
print(df['Owned_Work_Phone'].value_counts())
print("\n")
```

```
# Check unique values and counts for 'Owned_Phone'
print("Unique values and counts for 'Owned_Phone':")
print(df['Owned_Phone'].value_counts())
print("\n")
```

```
# Check unique values and counts for 'Owned_Email'
print("Unique values and counts for 'Owned_Email':")
print(df['Owned_Email'].value_counts())
```

```
#Unique values and counts for 'Owned_Mobile_Phone':
#1      25128
#Name: Owned_Mobile_Phone, dtype: int64
```

```
#Unique values and counts for 'Owned_Work_Phone':
#0      18249
#1       6879
#Name: Owned_Work_Phone, dtype: int64
```

```
#Unique values and counts for 'Owned_Phone':
#0      17772
#1       7356
#Name: Owned_Phone, dtype: int64
```

```
#Unique values and counts for 'Owned_Email':
#0      22598
#1       2530
#Name: Owned_Email, dtype: int64
```

```
#It seems like for the attribute 'Owned_Mobile_Phone', all the values are '1'.
#The column is constant, there are no varied values, so, we will drop it.
```

```
df.drop('Owned_Mobile_Phone', axis=1, inplace=True)
```

```
##%
```

```
# Also, the 'Applicant_ID' should be having separate values for all rows. So, it
wouldn't be adding much
#values to the analysis, so we will drop it.
```

```
df.drop('Applicant_ID', axis=1, inplace=True)
```

```

print(df.head())

print(df.info())

###

## Creating a boxplot for all numerical attributes to check for outliers

# Define numerical columns to plot
numerical_columns = ['Total_Income', 'Total_Family_Members', 'Applicant_Age',
                     'Years_of_Working', 'Total_Bad_Debt', 'Total_Good_Debt']

# Create separate box plots for each numerical column
sns.set(style="whitegrid")
for col in numerical_columns:
    plt.figure(figsize=(8, 5))
    sns.boxplot(x=df[col], color='skyblue')
    plt.title(f'Box Plot for {col}')
    plt.xlabel('')
    plt.show()

# Most of our plots still have decent amount of values over the 75th percentile.
# Using 1th & 99th percentile as our benchmark for removing our outliers

###

# Removing outliers

# Defining numerical columns to filter
numerical_columns = ['Total_Income', 'Total_Family_Members', 'Years_of_Working',
                     'Total_Bad_Debt']

# Calculating 1th and 99th percentiles for each column
percentiles = df[numerical_columns].quantile([0.01, 0.99])

# Filtering DataFrame to remove outliers
filtered_df = df.copy() # Create a copy of the DataFrame
for col in numerical_columns:
    lower_bound = percentiles.loc[0.01, col]
    upper_bound = percentiles.loc[0.99, col]
    filtered_df = filtered_df[(filtered_df[col] >= lower_bound) & (filtered_df[col]
<= upper_bound)]

# Display
print("Summary Statistics Before Outlier Removal:")
print(df[numerical_columns].describe())
print("\nSummary Statistics After Outlier Removal:")
print(filtered_df[numerical_columns].describe())

```

#Summary Statistics After Outlier Removal:

#	Total_Income	Total_Family_Members	Years_of_Working	Total_Bad_Debt
#count	24130.000000	24130.000000	24130.000000	24130.000000
#mean	190360.741442	2.284252	7.421508	0.213717
#std	85019.559173	0.904108	5.856716	0.735050
#min	63000.000000	1.000000	1.000000	0.000000

```
#25%    135000.000000    2.000000    3.000000    0.000000
#50%    180000.000000    2.000000    6.000000    0.000000
#75%    225000.000000    3.000000   10.000000    0.000000
#max    585000.000000    5.000000   30.000000    6.000000
```

```
####
```

```
print(filtered_df.info())
print(filtered_df.shape)
```

```
####
```

```
#Converting the categorical features into numeric
```

```
categorical_features = ['Applicant_Gender',
'Owned_Car', 'Owned_Realty', 'Income_Type', 'Education_Type',
'Family_Status', 'Owned_Work_Phone', 'Owned_Phone', 'Owned_Email', 'Housing_Type',
'Job_Category']
```

```
# converting object-type categorical features to dummy variables
filtered_df = pd.get_dummies(filtered_df, columns=categorical_features,
drop_first=True)
```

```
print(filtered_df.head())
print(filtered_df.info())
```

```
####
```

```
#checking for under/over sampling
```

```
status_distribution = filtered_df['Status'].value_counts()
print("Class Distribution:")
print(status_distribution)
```

```
# plotting distribution of status
status_distribution.plot(kind='bar', rot=0)
plt.title("Class Distribution of Status")
plt.xlabel("Status")
plt.ylabel("Count")
plt.show()
```

```
# calculating percentages of status distribution (%)
class_proportions = status_distribution / len(filtered_df) * 100
print("\nClass Proportions (%):")
print(class_proportions)
```

```
# observation - the data of people whose credit got approved is clearly more,
# meanwhile the folks who didn't get their credit approved have quite less
representation in the dataset.
```

```
#Class Proportions (%):
#1    99.772068
#0     0.227932
```

```
####
```

```
### have to use a way to prevent oversampling while splitting/training/testing
```

```

### Might have to check the collinearity of attributes

###

# Separating features (X) and target variable (y)
X = filtered_df.drop('Status', axis=1) # Features
y = filtered_df['Status']             # Target variable

###

from imblearn.pipeline import Pipeline as ImbPipeline
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, roc_curve, roc_auc_score, confusion_matrix
import matplotlib.pyplot as plt

# Splitting the data into training and test sets (stratified sampling for
imbalance)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
stratify=y, random_state=42)

# Applying SMOTE to the training data only
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

###

##### MODEL - I ##### Logistic Regression

# Printing value counts of target variables after oversampling
print("After oversampling:")
print("Train set - Class 1:", sum(y_train_resampled == 1), "Class 0:",
sum(y_train_resampled == 0))

# Defining the parameter grid for logistic regression
log_grid = {
    'clf__C': np.logspace(-2, 5, 21), # Regularization parameter
    'clf__penalty': ['l1', 'l2'],      # Type of penalty compatible with LBFGS
    'solver'
        'clf__solver': ['liblinear'], # Solver algorithm
        'clf__class_weight': ['balanced', None], # Class weight for imbalance
        'clf__max_iter': [100, 200, 500] # Maximum number of iterations
    }

# Defining the pipeline with standard scaling and logistic regression
estimator = ImbPipeline([
    ('scale', StandardScaler()),
    ('clf', LogisticRegression(random_state=10))
])

```

```

# Creating a GridSearchCV object with the pipeline
logistic = GridSearchCV(estimator=estimator,
                        param_grid=log_grid,
                        cv=5,
                        scoring='f1',
                        n_jobs=-1)

# Fitting the GridSearchCV object on the resampled training data
logistic.fit(X_train_resampled, y_train_resampled)

# Printing the best hyperparameters found
print("Best Hyperparameters:", logistic.best_params_)

# Performing evaluation on test set
y_pred = logistic.predict(X_test)
f1 = f1_score(y_test, y_pred)
print("F1-score on test set:", f1)

# Calculating other metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score

precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)

print("Precision on test set:", precision)
print("Recall on test set:", recall)
print("Accuracy on test set:", accuracy)

# Plotting ROC curve
fpr, tpr, _ = roc_curve(y_test, logistic.predict_proba(X_test)[:, 1])
roc_auc = roc_auc_score(y_test, logistic.predict_proba(X_test)[:, 1])

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

# confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

# finding feature importance
coef = logistic.best_estimator_.named_steps['clf'].coef_[0]
feature_names = X.columns
feature_importance = pd.DataFrame({'Feature': feature_names, 'Importance': coef})
top_10_features = feature_importance.sort_values(by='Importance',
ascending=False).head(10)
print("Top 10 Important Features:")
print(top_10_features)

```



```

import seaborn as sns

# plotting feature importance
plt.figure(figsize=(10, 8))
sns.barplot(x='Importance', y='Feature', data=top_10_features)
plt.title('Top 10 Important Features')
plt.show()

# =====
# After oversampling:
# Train set - Class 1: 19260 Class 0: 19260
# Best Hyperparameters: {'clf__C': 44668.35921509626, 'clf__class_weight':
# 'balanced', 'clf__max_iter': 100, 'clf__penalty': 'l2', 'clf__solver': 'liblinear'}
# F1-score on test set: 1.0
# Precision on test set: 1.0
# Recall on test set: 1.0
# Accuracy on test set: 1.0
# Confusion Matrix:
# [[ 11   0]
#  [  0 4815]]
# Top 10 Important Features:
#
#                                     Feature Importance
# 6                                     Total_Good_Debt 113.980978
# 17 Education_Type_Secondary / secondary special ... 12.361044
# 14 Education_Type_Higher education ... 10.073297
# 33 Job_Category_Unskilled Labor Roles 7.066006
# 30 Job_Category_Professional Roles 5.504682
# 15 Education_Type_Incomplete higher ... 5.132137
# 31 Job_Category_Service and Support Roles 5.082376
# 32 Job_Category_Skilled Labor Roles 4.380168
# 2 Total_Family_Members 3.470471
# 21 Family_Status_Widow ... 2.522999
#
# =====

#%#%

#%#%

#### MODEL - II ##### RANDOM FOREST

from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import f1_score, roc_curve, roc_auc_score, confusion_matrix,
classification_report, precision_recall_curve

#Defining the parameter grid for random forest
rf_grid = {'n_estimators': np.linspace(100, 1000, 10, dtype = int),
           'max_depth': [None, 10, 20, 30],
           'min_samples_split': [2, 5, 10],
           'max_features': ['sqrt', 'log2', None]}

```

```
RF = GridSearchCV(RandomForestClassifier(min_samples_leaf = 10, random_state = 10,
max_features = 'sqrt'),
                  param_grid = rf_grid, cv = 5, n_jobs = -1, scoring = 'f1')
RF.fit(X_train_resampled, y_train_resampled)
```

```
###
```

```
# Printing the best hyperparameters found
print("Best Hyperparameters:", RF.best_params_)
```

```
# Getting the best estimator from GridSearchCV
best_rf_model = RF.best_estimator_
```

```
# Extracting feature importances from the best model
importances = best_rf_model.feature_importances_
```

```
# Getting the names of the features/columns
feature_names = X_train.columns
```

```
# Create a DataFrame to store feature names and their importances
feature_importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
importances})
```

```
# Sorting the DataFrame by importance in descending order
feature_importance_df = feature_importance_df.sort_values(by='Importance',
ascending=False)
```

```
# Printing the top 10 important attributes
print("Top 10 Important Attributes:")
print(feature_importance_df.head(10))
```

```
# Plotting the top 10 important features
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df.head(10))
plt.title('Top 10 Important Features')
plt.show()
```

```
# =====
# #Best Hyperparameters: {'max_depth': 20, 'max_features': 'sqrt',
# 'min_samples_split': 2, 'n_estimators': 200}
# #Top 10 Important Attributes:
#
#           Feature  Importance
# 6      Total_Good_Debt    0.355194
# 5      Total_Bad_Debt    0.341899
# 9      Owned_Realty_1    0.029581
# 13 Income_Type_Working    0.027199
# 8      Owned_Car_1      0.026915
# 31      Job_Category_Service and Support Roles    0.024955
# 33      Job_Category_Unskilled Labor Roles    0.021827
# 23      Owned_Phone_1    0.020966
# 30      Job_Category_Professional Roles    0.020873
# 22      Owned_Work_Phone_1    0.014933
# =====
```

```

###
# Performing evaluation on holdout set
from sklearn.metrics import f1_score
print(f1_score(y_test, RF.predict(X_test)))

###
# ROC Curves
from sklearn.metrics import roc_curve, roc_auc_score
fpr_rf, tpr_rf, _ = roc_curve(y_train, RF.predict_proba(X_train)[:,-1])

###
plt.figure(figsize=(10, 6))
plt.plot(fpr_rf, tpr_rf, label=f'Random Forest (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")
plt.show()

# Generating Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, RF.predict_proba(X_test)[:,-1])
plt.figure(figsize=(10, 6))
plt.plot(recall, precision, marker='.', label='Random Forest')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend()
plt.show()

# Printing classification report
print(classification_report(y_test, y_pred))

# Printing AUC values
print("AUC Score on test set:", roc_auc)

# Confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=[1, 0])
print("Confusion Matrix:")
print(cm)
# =====
#
#
#           precision    recall  f1-score   support
#
#    0           1.00        1.00        1.00         11
#    1           1.00        1.00        1.00        4815
#
#   accuracy                1.00         4826
#   macro avg           1.00        1.00        1.00         4826
#   weighted avg          1.00        1.00        1.00         4826
#
# AUC Score on test set: 1.0
# Confusion Matrix:
# [[4815    0]

```

```
# [ 0 11]
```

```
#
```

```
# =====
```